

Project:

This project was to implement gradient checking to ensure that the implementation of backpropagation is correct.

1-dimensional gradient checking

```
def forward_propagation(x, theta):
```

```
    """
```

```
        Implement the linear forward propagation (compute J)
```

```
(J(theta) = theta * x)
```

```
    Arguments:
```

```
    x -- a real-valued input
```

```
    theta -- our parameter, a real number as well
```

```
    Returns:
```

```
    J -- the value of function J, computed using the formula J(theta) = theta * x
```

```
    """
```

```
    J = theta*x
```

```
    return J
```

```
--
```

```
def gradient_check(x, theta, epsilon = 1e-7):
```

```
    """
```

```
    Implement the backward propagation
```

```
    Arguments:
```

```
    x -- a real-valued input
```

```
    theta -- our parameter, a real number as well
```

```
    epsilon -- tiny shift to the input to compute approximated gradient with  
    formula(1)
```

```
    Returns:
```

```
    difference -- difference (2) between the approximated gradient and the  
    backward propagation gradient
```

```
    """
```

```
    # Compute gradapprox
```

```
    thetaplus = theta+epsilon                                # Step 1
```

```
    thetaminus = theta-epsilon                               # Step 2
```

```
    J_plus = forward_propagation(x,thetaplus)               # Step 3
```

```
    J_minus = forward_propagation(x,thetaminus)             # Step 4
```

```
    gradapprox = (J_plus-J_minus)/(2*epsilon)               # Step 5
```

```
    # Check if gradapprox is close enough to the output of backward_propagation()
```

```
grad = backward_propagation(x,theta)
```

```
numerator = np.linalg.norm(grad-gradapprox) # Step 1'
```

```
denominator = np.linalg.norm(grad)+np.linalg.norm(gradapprox)  
# Step 2'
```

```
difference = numerator/denominator # Step 3'
```

```
if difference < 1e-7:
```

```
    print ("The gradient is correct!")
```

```
else:
```

```
    print ("The gradient is wrong!")
```

```
return difference
```

N-dimensional gradient checking

```
def forward_propagation_n(X, Y, parameters):
```

```
    """
```

```
    Implements the forward propagation (and computes the cost)
```

```
    Arguments:
```

```
    X -- training set for m examples
```

```
    Y -- labels for m examples
```

parameters -- python dictionary containing your parameters "W1", "b1", "W2", "b2", "W3", "b3":

W1 -- weight matrix of shape (5, 4)

b1 -- bias vector of shape (5, 1)

W2 -- weight matrix of shape (3, 5)

b2 -- bias vector of shape (3, 1)

W3 -- weight matrix of shape (1, 3)

b3 -- bias vector of shape (1, 1)

Returns:

cost -- the cost function (logistic cost for one example)

"""

retrieve parameters

m = X.shape[1]

W1 = parameters["W1"]

b1 = parameters["b1"]

W2 = parameters["W2"]

b2 = parameters["b2"]

W3 = parameters["W3"]

b3 = parameters["b3"]

```
# LINEAR -> RELU -> LINEAR -> RELU -> LINEAR -> SIGMOID
```

```
Z1 = np.dot(W1, X) + b1
```

```
A1 = relu(Z1)
```

```
Z2 = np.dot(W2, A1) + b2
```

```
A2 = relu(Z2)
```

```
Z3 = np.dot(W3, A2) + b3
```

```
A3 = sigmoid(Z3)
```

```
# Cost
```

```
logprobs = np.multiply(-np.log(A3),Y) + np.multiply(-np.log(1 - A3), 1 - Y)
```

```
cost = 1./m * np.sum(logprobs)
```

```
cache = (Z1, A1, W1, b1, Z2, A2, W2, b2, Z3, A3, W3, b3)
```

```
return cost, cache
```

```
--
```

```
def backward_propagation_n(X, Y, cache):
```

```
    """
```

```
    Implement the backward propagation
```

```
    Arguments:
```

```
    X -- input datapoint, of shape (input size, 1)
```

Y -- true "label"

cache -- cache output from forward_propagation_n()

Returns:

gradients -- A dictionary with the gradients of the cost with respect to each parameter, activation and pre-activation variables.

"""

m = X.shape[1]

(Z1, A1, W1, b1, Z2, A2, W2, b2, Z3, A3, W3, b3) = cache

dZ3 = A3 - Y

dW3 = 1./m * np.dot(dZ3, A2.T)

db3 = 1./m * np.sum(dZ3, axis=1, keepdims = True)

dA2 = np.dot(W3.T, dZ3)

dZ2 = np.multiply(dA2, np.int64(A2 > 0))

dW2 = 1./m * np.dot(dZ2, A1.T)

db2 = 1./m * np.sum(dZ2, axis=1, keepdims = True)

dA1 = np.dot(W2.T, dZ2)

dZ1 = np.multiply(dA1, np.int64(A1 > 0))

```
dW1 = 1./m * np.dot(dZ1, X.T)
```

```
db1 = 1./m * np.sum(dZ1, axis=1, keepdims = True)
```

```
gradients = {"dZ3": dZ3, "dW3": dW3, "db3": db3,  
             "dA2": dA2, "dZ2": dZ2, "dW2": dW2, "db2": db2,  
             "dA1": dA1, "dZ1": dZ1, "dW1": dW1, "db1": db1}
```

```
return gradients
```

```
--
```

```
def gradient_check_n(parameters, gradients, X, Y, epsilon = 1e-7):
```

```
    """
```

Checks if backward_propagation_n computes correctly the gradient of the cost output by forward_propagation_n

Arguments:

parameters -- python dictionary containing your parameters "W1", "b1", "W2", "b2", "W3", "b3":

grad -- output of backward_propagation_n, contains gradients of the cost with respect to the parameters.

x -- input datapoint, of shape (input size, 1)

y -- true "label"

epsilon -- tiny shift to the input to compute approximated gradient with formula(1)

Returns:

difference -- difference (2) between the approximated gradient and the backward propagation gradient

"""

Set-up variables

parameters_values, _ = dictionary_to_vector(parameters)

grad = gradients_to_vector(gradients)

num_parameters = parameters_values.shape[0]

J_plus = np.zeros((num_parameters, 1))

J_minus = np.zeros((num_parameters, 1))

gradapprox = np.zeros((num_parameters, 1))

Compute gradapprox

for i in range(num_parameters):

Compute J_plus[i]. Inputs: "parameters_values, epsilon". Output = "J_plus[i]".

"_" is used because the function you have to outputs two parameters but we only care about the first one


```

    thetaplus = np.copy(parameters_values)                                # Step 1

    thetaplus[i][0] = thetaplus[i][0]+epsilon                            # Step 2

    J_plus[i], _ = forward_propagation_n(X, Y, vector_to_dictionary(thetaplus))
# Step 3

    # Compute J_minus[i]. Inputs: "parameters_values, epsilon". Output =
    "J_minus[i]".

    thetaminus = np.copy(parameters_values)                                # Step 1

    thetaminus[i][0] = thetaplus[i][0]-epsilon                            # Step 2

    J_minus[i], _ = forward_propagation_n(X, Y,
vector_to_dictionary(thetaminus))                                # Step 3

    # Compute gradapprox[i]

    gradapprox[i] = (J_plus[i]-J_minus[i])/(2*epsilon)

    # Compare gradapprox to backward propagation gradients by computing
    difference.

    numerator = np.linalg.norm(grad-gradapprox)                            # Step 1'

    denominator = np.linalg.norm(grad)+np.linalg.norm(gradapprox)
# Step 2'

    difference = numerator/denominator                                    # Step 3'

    if difference > 2e-7:

        print ("\033[93m" + "There is a mistake in the backward propagation!")
        difference = " + str(difference) + "\033[0m")

```

else:

```
    print ("\033[92m" + "Your backward propagation works perfectly fine!"  
difference = " + str(difference) + "\033[0m")
```

return difference